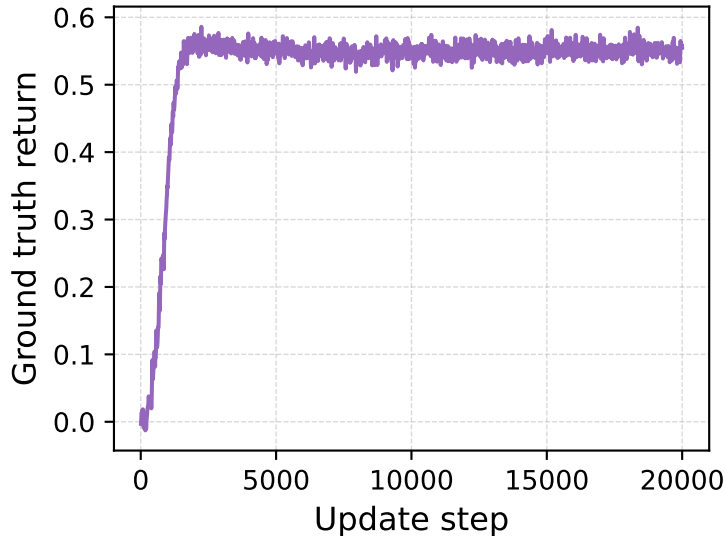


Learning curves at 10.0% oversight



Gradient-routed MoE